

Arvindh Arun

[arvindh75.github.io](https://github.com/arvindh75) [@ iiith@arvindharun.com](mailto:iiith@arvindharun.com) github.com/arvindh75 [Google Scholar](https://scholar.google.com/citations?user=arvindh75)
[in linkedin.com/in/arvindh75](https://www.linkedin.com/in/arvindh75) twitter.com/arvindh__a

Education

Current Aug 2024	ELLIS and IMPRS-IS PhD (ELLIS unit, Stuttgart & University of Edinburgh) PhD in Artificial Intelligence ‣ Advised by Prof. Steffen Staab and Prof. Antonio Vergari ‣ Training AI for Decision Making, LLM Evaluations, and Reasoning with Structured data	Germany
May 2024 Aug 2019	International Institute of Information Technology (IIIT H) B.Tech. + MS by Research, Computer Science & Engineering ‣ Cumulative GPA - 9.07/10.00 ‣ Dean's Academic List - 2020-21 M (Top 10%), 2021-22 M (Top 15%), 2022-23 S (Top 5%) ‣ Dean's Research List - 2022-23 (Top 5%) ‣ Represented the College Football Team, 5x medals in Track & Field, 3x medals in Inter-house football	Hyderabad, India

Experience

Sep 2026 May 2026	Sakana AI <i>Research Scientist Intern</i> ‣ Worked on pretraining analysis of Looped Language Models.	Tokyo, Japan
Aug 2024 Jan 2024	Google DeepMind <i>Student Research Collaborator</i> ‣ Worked on leveraging topological information from user interactions to gather insight into the quality of user-generated content. ‣ Extended Vietoris Rips Complex and Persistent Images to post profiling on social networks. ‣ Introduced a new benchmark to evaluate controversy detection models mimicking the real-world data distribution.	Remote
Jul 2023 May 2023	JPMorgan Chase Corporate Investment Banking <i>Applied AI/ML Intern</i> ‣ Worked on mining interesting insights from transactional databases for Root Cause Analysis. ‣ Utilized streaming frameworks such as PySpark to optimize the workflow. ‣ Proposed a new framework utilizing unsupervised GNNs inspired by the CTR prediction domain. ‣ Received a Pre-placement Full-time offer for my contributions.	Bangalore, India
Aug 2022 Jun 2022	The Chinese University of Hong Kong Department of Systems Engineering <i>Research Intern / Advisor: Prof. Hong Cheng</i> ‣ Worked on the domain of Clustering/Community detection in Dynamic social networks and studying their properties. ‣ Focused on minimizing the time complexity and the space complexity of the existing real-time solutions. ‣ Explored various state-of-the-art solutions and built a visualization tool to study their efficacy.	Hong Kong
May 2024 May 2021	Precog Research Group [🌐] <i>Undergraduate Researcher / Advisor: Prof. Ponnurangam Kumaraguru "PK"</i> ‣ Analyzing social networks and other large networks primarily using their network information. ‣ Georgia Institute of Technology, Atlanta Collaboration with Prof. Polo Chau on analyzing the Quantitative and Qualitative impact of fraudulent behavior on social networks and their identification. ‣ Adobe Research, Bangalore Collaboration with Dr. Ramasuri Narayanam on exploring the landscape of group fairness in unsupervised contrastive learning based GNNs. ‣ Indiana University, Bloomington Collaboration with Prof. Jisun An on analyzing the changes to the hate speech dynamics after the liberalization of social media platforms.	Hyderabad, India

Publications

- [6] **FutureSim: Replaying World Events to Evaluate Adaptive Agents** [PDF] [🔗]
Shashwat Goel*, Nikhil Chandak*, Arvindh Arun*, Ameya Prabhu, Steffen Staab, Moritz Hardt, Maksym Andriushchenko, Jonas Geiping (* = Equal Contribution)
40th Annual Conference on Neural Information Processing Systems (Evaluations & Datasets Track)
Best Paper *Forecasting as a New Frontier of Intelligence at ICML 2026* [NeurIPS '26]
- [5] **The Illusion of Diminishing Returns: Measuring Long Horizon Execution in LLMs** [PDF] [🔗]
Akshit Sinha*, Arvindh Arun*, Shashwat Goel*, Steffen Staab, Jonas Geiping (* = Equal Contribution)
14th International Conference on Learning Representations (ICLR) [ICLR '26]
- [4] **SEMMA: A Semantic Aware Knowledge Graph Foundation Model** [PDF] [🔗]
Arvindh Arun, Sumit Kumar, Mojtaba Nayyeri, Bo Xiong, Ponnurangam Kumaraguru, Antonio Vergari and Steffen Staab
30th Conference on Empirical Methods in Natural Language Processing (EMNLP) [EMNLP '25]
- [3] **A Cognac shot to forget bad memories: Corrective Unlearning in GNNs** [PDF] [🔗]
Varshita Kolipaka, Akshit Sinha, Debangan Mishra, Sumit Kumar, Arvindh Arun*, Shashwat Goel*, and Ponnurangam Kumaraguru (* = Equally Advising)
42nd International Conference on Machine Learning (ICML) [ICML '25]
- [2] **CAFIN: Centrality Aware Fairness inducing IN-processing for Unsupervised Representation Learning on Graphs** [PDF] [🔗]
Arvindh Arun, Aakash Aanegola, Amul Agrawal, Ramasuri Narayanam, and Ponnurangam Kumaraguru
Oral 26th European Conference on Artificial Intelligence (ECAI) [ECAI '23]
- [1] **Erasing Labor with Labor: Dark Patterns and Lockstep Behaviors on the Google Play Store** [PDF]
Ashwin Singh, Arvindh Arun, Pulak Malhotra, Pooja Desur, Ayushi Jain, Duen Horng Chau, and Ponnurangam Kumaraguru
33rd ACM Conference on Hypertext and Social Media (ACM HT)
Media Work covered by [Montreal AI Ethics Institute](#) [ACM HT '22]

Workshop Publications

- [W6] **Spotlight What Breaks Depth Extrapolation in Looped Language Models?**
Arvindh Arun*, Tianyu Zhao, Kai Arulkumaran
Linguistic Principles for Foundation Models at NeurIPS 2026 [LP4M @ NeurIPS '26]
- [W5] **Ex-GraphRAG: Interpretable Evidence Routing for Graph-Augmented LLMs**
Yoav Kor Sade*, Arvindh Arun*, Rishi Puri, Steffen Staab, and Maya Bechler-Speicher (* = Equal Contribution)
4th Towards Knowledgeable Foundation Models at ACL 2026 [KnowFM @ ACL '26]
- [W4] **Topo Goes Political: TDA-Based Controversy Detection in Imbalanced Reddit Political Data** [PDF]
Arvindh Arun*, Karuna K Chandra*, Akshit Sinha, Balakumar Velayutham, Jashn Arora, Manish Jain, and Ponnurangam Kumaraguru (* = Equal Contribution)
Best Paper 5th International Workshop on Computational Methods for Online Discourse Analysis @ WWW 2025 [BeyondFacts @ WWW '25]
- [W3] **Sanity Checks for Evaluating Graph Unlearning** [PDF]
Varshita Kolipaka, Akshit Sinha, Debangan Mishra, Sumit Kumar, Arvindh Arun, Shashwat Goel, and Ponnurangam Kumaraguru
3rd Conference on Lifelong Learning Agents Workshop Track (CoLLAS) 2024 [Workshop @ CoLLAS '24]
- [W2] **X-posing Free Speech: Examining the Impact of Moderation Relaxation on Online Social Networks** [PDF]
Arvindh Arun*, Saurav Chhatani*, Jisun An, and Ponnurangam Kumaraguru (* = Equal Contribution)
8th Workshop on Online Abuse and Harms (WOAH) at NAACL 2024 [WOAH @ NAACL '24]
- [W1] **Ensembling Pre-Trained Language Models with Feature Engineering** [PDF] [🔗]
T.H. Arjun, Arvindh Arun, and Ponnurangam Kumaraguru
GermEval 2021 Shared Task on the Identification of Toxic, Engaging, and Fact-Claiming Comments [GermEval '21]

Honours, Awards and Test Scores

Y Combinator AI Startup School, 2025 [🌐] Accepted to attend a funded 2-day AI startup school in the Bay Area.

Google Research Week, 2024 [🌐] Accepted to attend a funded 3-day ML research event, hosted by Google Research India.

Microsoft Research Travel Grant, 2023 [🌐] Awarded \$2500 for attending ECAI'23 held in Krakow, Poland.

Stellantis NV Hackathon, 2022 Finished as a runner-up in both rounds, winning a cash prize worth 125,000 INR.

Facebook Hackercup, 2022 Reached Round 2, finished Top 10% globally.

GermEval, 2021 Ranked 6th globally on cumulative task rank.

PicoCTF by CMU, 2021 Finished 102nd, Top 2% globally.

HowzHack (DreamVu), 2019 Finished 2nd for the problem statement exploring the applications of computer vision for Football's Video Assistant Referees (VAR).

UGEE, 2019 20th rank in India.

NextGenius Scholarship, 2019 Finalist (withdrew) selected for fully funded Undergraduate programme at University Of The South, Tennessee, USA.